

FraudOps-Bench: Process-Level, Tool-Grounded Evaluation of LLM Agents for Card-Not-Present Fraud Investigation, and a Calibration Pitfall in Selective Prediction with LLM Confidence Scores

AYUSHI AMBILKAR¹ AND PRAMEGH UIKEY²

¹Independent Researcher, India (e-mail: ambilkarayushi@gmail.com)

²Independent Researcher, India (e-mail: uikeyprameghaa@gmail.com)

Corresponding author: Ayushi Ambilkar (e-mail: ambilkarayushi@gmail.com).

ABSTRACT We study whether a large language model (LLM) agent can emulate a fraud analyst’s card-not-present review workflow: gathering evidence via tools, reasoning under an explicit standard operating procedure (SOP), and producing an auditable, evidence-cited disposition. FraudOps-Bench provides a six-tool investigation environment over IEEE Computational Intelligence Society (IEEE-CIS) transaction data, a LangGraph-based agent, two further LLM arms (a zero-evidence control and a single-shot baseline), and a classical gradient-boosted-tree baseline, evaluated under a selective-prediction framework that escalates low-confidence cases to a human. An initial 120-case calibration made the agent appear competitive with classical machine learning (ML), but the result did not survive a fresh 300-case holdout: the calibration procedure had overfit small-sample noise in the escalation threshold. We diagnose the cause—discretized, round-number-anchored LLM probability outputs interacting with a narrow band selected on too little data—fix it with a lower-confidence-bound selection criterion, and validate the fix on a second, disjoint holdout. Under corrected calibration the agent’s accuracy claim holds (96.1%), but its coverage collapses to 17% of the queue, and the classical baseline dominates the full risk-coverage curve (area under the risk-coverage curve, AURC, 0.034 vs. 0.108–0.112). Rerunning the identical protocol on a second model, GPT-5.6 Terra, reproduces the qualitative finding (AURC 0.185–0.205) with an even more extreme high-precision/low-coverage tradeoff. We also report two-method faithfulness scoring (~99% evidence-attribution accuracy across all four evidence-bearing arms), a structured error taxonomy, and a self-consistency ablation. The calibration failure mode is a transferable lesson for selective prediction with LLM confidence scores on small calibration sets.

INDEX TERMS agentic artificial intelligence, benchmark evaluation, fraud detection, large language models, model calibration, risk-coverage analysis, selective prediction, tool-augmented LLMs

I. INTRODUCTION

Fraud remains a large and continuous operational burden on financial institutions. The Association of Certified Fraud Examiners’ 2024 *Report to the Nations* documents thousands of occupational-fraud cases spanning 138 countries and territories [1], and UK Finance’s own H1 2025 figures report 2.09 million confirmed fraud cases and £629.3 million stolen

in just six months [2]—a reminder that the review queues behind these numbers are large, continuous, and staffed by human analysts working under time pressure.

Large language models (LLMs) are an obvious candidate to help with this queue, and a growing body of recent work has tested whether an LLM can classify a transaction as fraudulent as well as a classical machine learning

(ML) model trained on standard public benchmarks such as the IEEE Computational Intelligence Society’s (IEEE-CIS) fraud dataset [23]. That question is close to saturated: public fraud datasets are anonymized and feature-engineered for tabular models rather than investigative reasoning, and the strongest recent public-data-only result along these lines, FinFRE-RAG [4], still frames the task as transaction-level risk scoring rather than the workflow a human fraud analyst actually performs.

That workflow looks different from classification. A real analyst does not receive a fully assembled feature vector; they receive an alert, then decide what evidence to pull—card history, device history, prior velocity, identity-match signals—before reasoning under an institution’s standard operating procedure (SOP) toward a disposition: approve, escalate, or reject, with a citable rationale. Emulating that *process*, not just its final label, is the harder and less-studied problem. The closest published attempt is the FAA framework [5], which wraps GPT-4o’s Assistants application programming interface (API) around a multi-agent investigation pipeline and reports a very strong F1 score (the harmonic mean of precision and recall) of 0.98–0.99 on synthetic Sparkov/CCTD data—but on a limited, mostly synthetic sample, without a process-level evaluation of when the system should defer to a human rather than commit to a label. The nearest behavioral analogues to a tool-grounded, policy-constrained investigation agent come not from fraud but from security operations: CORTEX [6] and SIABench [7] both evaluate multi-agent, tool-using triage over auditable evidence in a security operations center (SOC) setting, but on non-public data and in a different operational domain. General-purpose agent benchmarks—SOP-Bench [8], τ -bench [9], IntellAgent [10]—supply the methodological template for evaluating long-horizon, policy-heavy, tool-using agents, but none target fraud operations specifically.

No existing public work, to our knowledge, combines (a) public fraud data, (b) a tool-grounded agent that gathers its own evidence rather than receiving it pre-assembled, and (c) a process-level, auditable evaluation that asks not just “was the label right” but “should this system have committed to a label at all.” That is the gap FraudOps-Bench targets.

Building and evaluating such a system surfaced a second, independently interesting result. Selective prediction—escalating low-confidence cases to a human rather than forcing a decision—is the natural evaluation frame for this kind of agent, since a fraud-review system that is wrong with full confidence is far more costly than one that is honest about uncertainty. Doing this well requires calibrating a decision threshold against a labeled set, and we found, by first getting it wrong, that LLM-generated confidence scores are not the well-behaved continuous outputs classical calibration methods assume: they cluster heavily on round, prompt-anchored values, and a threshold selected against a small calibration set can badly overfit that clustering in a way that only shows up on a much larger, independent holdout. We diagnose this failure precisely, fix it, and validate the fix on a second,

disjoint holdout set—and we report the failure mode itself as a first-class contribution, since it is a risk any selective-prediction pipeline built on LLM confidence scores over a modestly sized calibration set is likely to share.

Contributions:

- 1) FraudOps-Bench: a public-data (IEEE-CIS), tool-grounded, LangGraph-based fraud-investigation environment with an explicit SOP, four comparison arms, and a disciplined dev/calibration/holdout evaluation protocol with a frozen-methodology audit mechanism preventing evaluation leakage.
- 2) A documented case study of `calibrate_band()` overfitting on a small calibration set—diagnosed via probability-output discretization, fixed with a lower-confidence-bound + minimum-sample-size selection rule, and validated end-to-end on a second, never-touched holdout set.
- 3) Two independent faithfulness-evaluation methods (deterministic numeric cross-referencing and an LLM-judge pass) plus a structured error taxonomy, going beyond outcome accuracy to ask whether the agent’s cited reasoning is actually grounded in the evidence it retrieved.
- 4) An ablation isolating the contribution of a self-consistency mechanism to the agent’s calibrated probability estimates.
- 5) A cross-model replication check: the identical protocol—own from-scratch calibration, the corrected lower-confidence-bound (LCB) selection rule, a single-use holdout—applied to a second model, GPT-5.6 Terra (OpenAI), testing whether the calibration-discretization risk and the classical-ML-dominates-in-practice finding are Claude-specific or general.

The rest of the paper proceeds as follows. Section II positions this work against prior fraud-LLM and agent-benchmark literature. Section III describes FraudOps-Bench’s design. Section IV reports a two-method faithfulness evaluation of the LLM arms’ cited evidence. Section V is the central empirical contribution: the calibration failure, its diagnosis, fix, validation, and a replication check on a second model. Section VI reports a self-consistency ablation. Section VII reports cost and latency. Section VIII discusses implications, Section IX states limitations directly, and Section X concludes.

II. RELATED WORK

Fraud-specific LLM work. The Fraud Dataset Benchmark (FDB) [3] established the standard public-data infrastructure this line of work builds on, standardizing IEEE-CIS to a 561,013/28,527 train/test split over 67 retained features and explicitly cataloging the limitations of public fraud data—sparse identity fields, no case notes, no institutional history—that any LLM system built on it inherits. Within that constraint, FinFRE-RAG [4] is the cleanest recent demonstration that public-data-only LLM fraud scoring can be substantially improved by retrieval and feature reduction, but it remains

transaction-level risk scoring, not investigative workflow: the model receives a feature vector, not an alert it must build a case around.

The closest direct prior art to FraudOps-Bench is the FAA framework [5], which wraps GPT-4o’s Assistants API around a multi-agent pipeline—planning, evidence gathering, analysis, and report generation, with an optional vision agent—and reports very strong F1 (0.9801 on Sparkov, 0.99 on CCTD) across 500 evaluated transactions, with investigations averaging 5.5–7 steps. This is the paper closest in spirit to “an LLM agent investigates a fraud case,” but it evaluates on a limited, mostly synthetic sample and does not report a process-level or selective-prediction evaluation—there is no notion, in that evaluation, of the system declining to commit to a disposition.

Behavioral analogues outside fraud. The strongest evidence that tool-grounded, multi-agent decomposition helps on auditable, high-stakes investigation work comes not from fraud but from security operations. CORTEX [6] shows a multi-agent SOC-triage system improving actionable F1 from 0.66 to 0.78 and cutting false-positive rate from 24.9% to 14.2% over a single-agent baseline, on real production workflow traces. SIABench [7] evaluates frontier models on deep, tool-executing security-incident-analysis workflows and finds they still fail a large fraction of long-context scenarios (GPT-4o failed 11 of 25 in one setting)—direct evidence that process-level, tool-grounded evaluation surfaces failure modes that outcome-only scoring hides. Both operate on non-public data in a different operational domain from fraud, but both are closer behavioral analogues to what FraudOps-Bench measures than any fraud-specific paper we found.

Agent-evaluation methodology. SOP-Bench [8], τ -bench [9], and IntellAgent [10] supply the methodological template this paper adapts to fraud operations: long-horizon, policy-constrained, tool-using agent evaluation, as distinct from single-turn question answering. SOP-Bench in particular shows that larger tool registries can hurt task success (37% vs. 20.8% task-success rate in one ablation)—a caution directly relevant to FraudOps-Bench’s six-tool design. The Berkeley Function-Calling Leaderboard [11] evaluates the narrower function/tool-calling capability these systems depend on. FraudOps-Bench’s agentic arm is implemented on LangGraph [12], one of several general-purpose multi-agent orchestration frameworks in current use alongside AutoGen [13] and CrewAI [14]; none of these frameworks target fraud operations specifically, and none ship an evaluation protocol beyond their own tool-calling correctness.

LLM confidence calibration and selective prediction. Selective prediction—reject the input rather than force a decision when confidence is low—goes back to Chow’s rule [15] and was extended to deep networks by Geifman and El-Yaniv’s selective classification framework [16], both of which implicitly assume a well-behaved, continuous confidence signal. Whether LLM-elicited confidence scores meet that assumption is itself an active question: Tian et al. [17] show that prompting a reinforcement-learning-from-human-

feedback (RLHF)-tuned model to verbalize a probability can be better calibrated than its token-level likelihood, but Xiong et al.’s systematic evaluation [18] finds elicited confidence remains frequently overconfident and unstable across prompting strategies. Most directly relevant to Section V-D’s root-cause finding, Dai and Wang [19] report that verbalized LLM confidence concentrates on a small number of round-number values across six models and three datasets— independent confirmation that the discretization this paper diagnoses in `agentic_api`’s outputs is a general property of LLM confidence elicitation, not an artifact specific to this benchmark or model. Conformal prediction [20] offers an alternative, distribution-free route to calibrated uncertainty for LLM outputs, but typically targets prediction *sets* over a fixed label space rather than the single scalar probability this benchmark’s SOP requires the model to commit to as part of an auditable disposition; we adopt the lower-confidence-bound-based fix (Section V-E) as the more direct match to that constraint, not as a claim that it dominates conformal alternatives in general.

Adjacent public fraud datasets. BAF [21] and AML-world [22] are the strongest synthetic-from-real alternatives to IEEE-CIS for future extensions of this benchmark—BAF in particular would extend FraudOps-Bench to account-opening fraud, a structurally different investigation shape from the card-not-present setting studied here (Section X).

Positioning. Distinguishing this paper from the saturated “LLM vs. classical ML on IEEE-CIS” classification literature [3], [4] is deliberate: FraudOps-Bench’s contribution is not a new state-of-the-art classification number, but a process-level, tool-grounded evaluation protocol—including a documented selective-prediction failure mode and its fix—applied to a domain that, to our knowledge, no existing public benchmark combines with a genuinely tool-grounded, multi-step investigation agent.

III. FRAUDOPS-BENCH DESIGN

A. TASK AND DATA

The task is card-not-present transaction risk review over the IEEE-CIS Kaggle fraud data [23], the same underlying dataset that FDB [3] uses as its richest public-data testbed. FraudOps-Bench draws its own dev/calibration/holdout splits directly from the raw competition data (Section III-E) rather than reusing FDB’s fixed train/test split, since our evaluation needs disjoint labeled subsets for calibration and repeated holdout use rather than a single fixed split. Each case exposes a visible alert summary (amount, product code, card type, purchaser/recipient email domain, device type) and, for evidence-bearing arms, access to six tools mirroring what a real fraud analyst’s case-management tooling would expose: `get_transaction_details`, `get_card_history`, `get_email_domain_profile`, `get_device_history`, `get_velocity_summary`, `get_identity_match_summary`. Tool outputs are leakage-free by construction (only information available strictly before the current transaction’s timestamp).

Case sampling, and what it does and does not represent. Two sampling decisions shape every number in this paper and are stated here rather than left implicit. First, cases are drawn only from transactions carrying device information (a non-null DeviceType or DeviceInfo) alongside a non-null card identifier and amount, since an investigation arm cannot exercise `get_device_history` or `get_identity_match_summary` on a transaction with no identity record at all. That eligible pool is 140,855 of IEEE-CIS’s 590,540 transactions (23.9%), and is itself higher-risk than the dataset as a whole: 7.96% fraud versus 3.50% overall. Second, every split is drawn class-balanced—50% fraud by construction—so that a 120-case calibration set and a 300-case holdout contain enough positives to estimate per-band accuracy at all.

The consequence is that this benchmark measures *discrimination under a balanced queue*, not calibration to an operational base rate. Precision, coverage-weighted correctness, and the risk-coverage curves in Section V would all shift—precision most sharply—if the same arms were run against a queue at the 3.5% prevalence of the raw dataset, and no result here should be read as a deployment-realistic false-positive estimate (Section IX). The comparison *between* arms is unaffected, since all arms see the identical cases.

B. COMPARISON ARMS

TABLE 1. Comparison arms.

Arm	Description
direct_control	Sees only the visible alert summary; no tool access. Zero-evidence control.
linear_api	Single-shot completion given the full output of all six tools at once.
agentic_api	A LangGraph agent with explicit brain (LLM), orchestrator, tool nodes, structured memory, and a supervisor node that checks required-check completion before allowing finalization. Iterative: the agent chooses which tools to call and when.
classical_ml	HistGradientBoostingClassifier on leakage-free tabular features derived from the same underlying data (no LLM).

Fig. 1 shows agentic_api’s LangGraph architecture. Section V reports all four arms above as run on Claude Sonnet 5 [24], the paper’s primary model; Sections IV and VI cover the two evidence-bearing LLM arms. Section V-H additionally reports the identical linear/agentic flows run on a second model, GPT-5.6 Terra (OpenAI)—denoted linear_gpt/agentic_gpt—as a cross-model replication check on the calibration methodology and the practical-performance finding. The local-model variants (linear_local/agentic_local, built but not run—see Limitations) remain out of scope for this report.

C. IMPLEMENTATION DETAILS

LLM arms. The Claude arms run claude-sonnet-5 and the cross-model arms gpt-5.6-terra, both through their providers’ hosted APIs. Decoding parameters are left at provider

defaults throughout: neither model accepts a non-default temperature or top_p on these endpoints (both return HTTP 400), so no sampling-parameter tuning was performed on any arm—a constraint that happens to remove one degree of freedom from the comparison. Response caps were set per backend by piloting on the dev split until no response truncated mid-JSON: 8192 output tokens for Anthropic, 16,384 for OpenAI. Failed calls retry up to five times with exponential backoff (2 s base, 60 s cap); Section VII reports the observed failure rate.

classical_ml. A HistGradientBoostingClassifier (scikit-learn, class_weight="balanced", random_state=0, otherwise library defaults—no hyperparameter search was run) over 42 engineered features: prior counts and prior fraud rates per card, per device and per email domain; prior average amounts and amount-versus-group ratios per card and per device; a purchaser/recipient domain-match flag; aggregates over the C, D and M column families; and one-hot encodings of ProductCD, card4, card6 and DeviceType. Group aggregates are computed strictly over rows preceding the target transaction, matching the leakage-free ordering the tool layer uses. The model trains on the full 590,540-row competition table with every benchmark transaction identifier removed—the union of dev, calibration, holdout and holdout_v2, enforced in code rather than by convention—so no case any arm is scored on appears in its training data. A logistic-regression variant exists in the repository but is not reported here; all classical_ml numbers in this paper are the gradient-boosted model.

D. SELECTIVE PREDICTION AND THE SOP

Every LLM arm outputs a calibrated fraud_probability. Disposition (APPROVE/ESCALATE/REJECT) is a deterministic function of that probability against a per-arm band, not a free-form model choice—the standard selective-prediction (Chow’s rule [15]) risk-coverage framing: fix an acceptable error rate on decided cases, escalate the rest. The SOP explicitly instructs the model on a calibration scale (0.0–0.15 no signal, 0.35–0.65 genuinely mixed evidence, 0.85–1.0 strong convergent signal) and requires an explicit per-check verdict (protective/risk/neutral) for six required checks before a final probability, to counter base-rate-neglect and encourage genuine escalation on ambiguous cases rather than false confidence.

Metrics. Write $p_i \in [0, 1]$ for the probability an arm assigns to case i and $y_i \in \{0, 1\}$ for its label. Bands are symmetric about the neutral point, $(\ell, u) = (0.5 - \tau, 0.5 + \tau)$ for a half-width τ , and induce a selection indicator and a prediction

$$s_i(\tau) = \mathbf{1}[p_i \leq \ell \text{ or } p_i \geq u], \quad \hat{y}_i = \mathbf{1}[p_i \geq u], \quad (1)$$




FIGURE 1. agentic_api's LangGraph architecture (Table 1). The brain (LLM) and orchestrator iteratively select among six tools, writing structured verdicts to memory each round, until the brain signals it is ready to finalize; the supervisor node then checks that all six required checks were actually completed before allowing a disposition, routing back to the orchestrator rather than allowing an under-evidenced finalization if not. linear_api takes the single-shot path (alert directly to brain with all six tools' output pre-attached, no orchestrator loop); direct_control sees only the alert box.

where $s_i = 0$ means the case is escalated to a human rather than decided. Over n cases, coverage and selective risk are

$$c(\tau) = \frac{1}{n} \sum_{i=1}^n s_i(\tau), \quad R(\tau) = \frac{\sum_i s_i(\tau) \mathbf{1}[\hat{y}_i \neq y_i]}{\sum_i s_i(\tau)}, \quad (2)$$

and the tables below report selective accuracy $1 - R(\tau)$ under the heading "Acc." Because arms escalate at very different rates, we also report the coverage-weighted score

$$CW(\tau) = c(\tau) (1 - R(\tau)), \quad (3)$$

the fraction of the whole queue an arm resolves correctly. Sweeping τ over a grid traces the arm's full risk-coverage curve (Fig. 2), summarized by its area,

$$AURC = \int R dc, \quad (4)$$

estimated by the trapezoid rule over the realized $(c(\tau), R(\tau))$ pairs for $\tau \in \{0, 0.025, \dots, 0.5\}$, sorted by coverage. Lower AURC is better, and unlike any single operating point it does not depend on which band an arm happened to select.

E. EVALUATION PROTOCOL: DEV / CALIBRATION / HOLDOUT, WITH A FROZEN-METHODOLOGY GUARD

To prevent the eval-leakage failure mode common in iterative prompt engineering (tuning a prompt or threshold by repeatedly checking accuracy on the set being reported), FraudOps-Bench uses three disjoint splits:

- **dev** ($n = 50$, 25/25 fraud): free to use for iterative SOP/prompt development. Already spent this way; never reported as a final number.
- **calibration** ($n = 120$, 60/60 fraud): used only to select each arm's escalate-probability band via k -fold cross-validation.
- **holdout** ($n = 300$, 150/150 fraud, two independent draws used in this work—see Section V): single-use final reporting set.

A methodology-freeze mechanism (freeze_methodology.py) hashes every file that defines an arm's behavior (SOP text, model config, prompt-construction code, agent graph definition, calibration logic) into a manifest. The holdout-execution path refuses to run unless the current files match

the frozen hashes, with a loud, logged --force override available but never silent. This is the mechanism that makes the before/after account in Section V possible: it is verifiable, not merely asserted, that nothing was returned against the second holdout after the fix was frozen.

IV. FAITHFULNESS EVALUATION

Two independent methods, both described in docs/methodology_log.md's 2026-08-14 and 2026-08-18 entries, are used to check whether an arm's cited reasoning is actually grounded in the evidence it retrieved.

Method 1—deterministic numeric cross-referencing. Extracts every numeric claim from an arm's cited evidence (evidence_used, check_verdicts, risk_indicators, etc.) and checks it against the real tool-evidence values (including a percentage-form match for fractional fields). Zero additional API cost. Across dev, calibration, holdout, and holdout_v2, both LLM arms show 98.7–99.3% verified rates. Residual "unverified" claims are overwhelmingly the model correctly performing arithmetic (e.g. a time delta between two timestamps) that the ground-truth flattener does not itself compute—not fabrication (manually verified, see docs/methodology_log.md).

Method 2—LLM-judge semantic check. A judge based on Claude Haiku 4.5 (claude-haiku-4-5; 25-case samples per arm on holdout_v2) checks for semantic misattribution the numeric method cannot catch: a correct number attached to a false claim. Found genuine (if minor) errors in 6/24 (linear_api) and 8/24 (agentic_api) scored cases (one judge-response parse failure excluded per arm)—see docs/error_taxonomy.md for the full categorization, including one category (rounding-precision framing) that is arguably not a real error but a limitation of the judge's own literalness, noted for completeness.

Result: both arms are highly faithful to the evidence they actually retrieve. The accuracy differences reported in Section V are not explained by evidence fabrication.

V. THE CALIBRATION FAILURE, DIAGNOSIS, FIX, AND VALIDATION

This section is presented as a first-class result, not an appendix correction—the failure mode is general to selective

prediction with LLM-generated probabilities and worth other practitioners knowing about.

A. INITIAL RESULT (CALIBRATION SPLIT, $N = 120$)

Band selection via 5-fold cross-validated accuracy on the calibration split initially favored a very narrow escalate band for agentic_api, (0.45, 0.55), yielding 85.7% accuracy at 99% coverage—appearing competitive with classical_ml (86.0% accuracy, 95% coverage).

B. REPEAT-RUN VARIANCE (STOCHASTICITY CHECK)

A competing explanation for the holdout regression (Section V-C) is ordinary generation-time stochasticity rather than anything about the band. To size that, each API arm was re-run twice on the unchanged $n = 120$ calibration split (rep1, rep2, identical prompts and configuration). In reporting order this measurement precedes the diagnosis it helps rule out, but chronologically it followed the fix: the runs were made after calibrate_band() had been corrected, and all three runs per arm are scored under the corrected band (Section V-E) so that the three are comparable to each other.

TABLE 2. Repeat-run variance on the $n = 120$ calibration split.

Arm	Run accuracies	Mean	Std
linear_api	0.912, 0.905, 0.880	0.899	± 0.017
agentic_api	0.938, 0.969, 0.966	0.957	± 0.017

Run-to-run variance from stochasticity alone is small ($\sim \pm 1.7$ points for both arms)—much smaller than the sampling-noise confidence interval (CI) at $n = 120$ ($\sim \pm 6$ pt) or even $n = 300$ ($\sim \pm 4$ pt, Section V-F). Single runs are not being meaningfully misled by generation-time randomness; whatever explains the ~ 10 -point holdout regression in Section V-C has to be something other than ordinary model stochasticity. This measurement (referenced again in Section VI) is the closest available estimate of single-run noise for both LLM arms, and it is what rules out stochasticity as the explanation, leaving the selection procedure itself as the remaining candidate.

C. FIRST HOLDOUT ($N = 300$) AND THE REGRESSION

Under the frozen (0.45, 0.55) band, agentic_api scored only 75.9% accuracy on a fresh 300-case holdout—a ~ 10 -point drop. classical_ml held steady (87.3%).

D. ROOT CAUSE

agentic_api's fraud_probability outputs are not continuous: they cluster heavily on round 0.05-increment values, with 0.45 and 0.55—exactly the selected band's edges—the two single most common outputs (53/300 and 45/300 cases on the holdout, respectively). On the 120-case calibration split, only 12 and 13 cases landed at those exact values, and by

sampling luck those small groups skewed hard toward one class (75% non-fraud at 0.45, 61.5% fraud at 0.55), making the narrow band look excellent. On the larger holdout, the same exact values are close to coin flips (47.2% and 55.6% fraud respectively)—consistent, in fact, with what the SOP's own calibration-scale instructions say those values should mean (genuinely mixed evidence). The calibration procedure mistook small-sample noise for signal.

E. FIX

calibrate_band() was changed from selecting the narrowest band whose mean cross-validated accuracy cleared a target threshold, to selecting the narrowest band whose **LCB** (mean minus 1.645 standard errors across folds) clears the threshold, and additionally requiring a minimum total decided-case count across folds regardless of how good a band's point estimate looks. Re-run on the same calibration data, this correctly rejects the narrow band (LCB 0.794 vs. an 0.85 target) and selects (0.25, 0.75) instead (LCB 0.895).

F. VALIDATION ON A SECOND, INDEPENDENT HOLDOUT ($N = 300$)

A fresh, disjoint 300-case holdout (holdout_v2, excluded from the classical-model training set along with every prior split) was generated and run under the corrected, re-frozen methodology.

TABLE 3. holdout_v2 results ($n = 300$).

Arm	Band	Cov.	Acc.	Prec.	Rec.	F1	Cov.-wt.
linear_api	(.35,.65)	50.7%	0.914	0.917	0.946	0.931	46.3%
agentic_api	(.25,.75)	17.0%	0.961	0.944	1.000	0.971	16.3%
classical_ml	(.35,.65)	85.3%	0.910	0.876	0.971	0.921	77.7%
direct_control	(.4,.6)	0%	—	—	—	—	—

“Coverage-weighted” = coverage $\times$ accuracy, the fraction of the full queue resolved correctly—the fairer single-number comparison when arms escalate at very different rates. direct_control is shown at the default (0.4, 0.6) band; it decided no cases.

Table 4: same three arms, matched at four notable bands. Table 3 compares each arm at its own independently-selected operating point, which is the correct primary comparison but obscures how each arm performs off that point. Since disposition is a deterministic function of fraud_probability and the band, every arm's accuracy/coverage at *any* band is a free re-computation on the already-collected holdout_v2 probabilities (Section V-F's Fig. 2 curve)—no new API calls. Shown here for all three arms at the four bands discussed in this paper: the original overfit band, the common band linear_api/classical_ml were independently selected at, the wide band agentic_api was independently selected at, and the (0.2, 0.8) band linear_gpt was independently selected at (Section V-H). Each arm's own selected operating point (the one reported in Table 3) is **bolded**.

TABLE 4. All three arms at four notable bands.

Arm	(.45,.55)	(.35,.65)	(.25,.75)	(.2,.8)
linear_api	cov 96.7%, acc 0.752	cov 50.7%, acc 0.914	cov 37.3%, acc 0.929	cov 24.7%, acc 0.959
agentic_api	cov 99.7%, acc 0.759	cov 39.7%, acc 0.924	cov 17.0%, acc 0.961	cov 10.0%, acc 0.967
classical_ml	cov 95.3%, acc 0.888	cov 85.3%, acc 0.910	cov 73.7%, acc 0.937	cov 68.7%, acc 0.961

Where each method shines, read directly off this table:

- **classical_ml's advantage is holding accuracy while staying willing to decide.** At the tightest band it still resolves 95.3% of the queue at 88.8% accuracy—both other arms are 20+ points less accurate at that same near-total coverage (0.752, 0.759). This is the source of its risk-coverage-curve dominance (Section V-F, Fig. 2): it does not trade coverage for accuracy nearly as steeply as the LLM arms do.
- **agentic_api's advantage is ceiling accuracy at its most selective points.** At its selected (0.25, 0.75) band it reaches 0.961—edging out classical_ml's 0.937 at the same band—and at (0.2, 0.8) it reaches the highest accuracy in the table (0.967), but at less than a quarter of classical_ml's coverage at either band (17.0% vs. 73.7%; 10.0% vs. 68.7%). When agentic_api commits, its confidence is well-placed; it simply commits rarely.
- **linear_api never wins a cell outright** but is never far behind either, sitting between the other two on both axes at every band—consistent with its middling coverage-weighted score in Table 3.
- The practical reading: classical_ml is the stronger choice as a primary decision-maker across the board; agentic_api's comparative advantage is real but narrow—a high-precision second opinion on the small slice of cases it is willing to rule on, not a general-purpose replacement.

TABLE 5. Bootstrap 95% CIs (10,000 resamples, src/stats_analysis.py).

Arm	Accuracy	Coverage	Cov.-wt.
linear_api	0.915 [0.868, 0.957]	0.507 [0.450, 0.563]	0.463 [0.407, 0.520]
agentic_api	0.961 [0.900, 1.000]	0.170 [0.130, 0.213]	0.164 [0.123, 0.207]
classical_ml	0.910 [0.873, 0.944]	0.853 [0.813, 0.893]	0.777 [0.727, 0.823]

TABLE 6. Pairwise McNemar's tests (paired, restricted to cases both arms decided). *b* and *c* are the discordant-pair counts: cases the first arm got right and the second wrong, and the reverse.

Comparison	<i>n</i> (both)	<i>b</i>	<i>c</i>	<i>p</i>
linear_api vs. agentic_api	50	0	0	1.0000
linear_api vs. classical_ml	137	4	5	1.0000
agentic_api vs. classical_ml	47	0	2	0.5000

Because arms with very different coverage (17% vs. 85%) share few commonly-decided cases, these paired tests are underpowered by construction (see the statistics report in the supplementary material) and are reported for completeness rather than as the primary evidence—the bootstrap CIs and risk-coverage curves above are more informative when coverage differs this much.

G. INTERPRETATION

The agentic_api accuracy claim is real. It holds up on completely fresh data under a properly calibrated band—96.1%, its best result on any split—confirming the original problem was the band, not a capability regression. But the cost is now visible: at that accuracy, it escalates 83% of the queue. classical_ml remains the strongest *practical* performer by a wide margin, both by coverage-weighted correctness (77.7% vs. 16.3%) and by dominating the full risk-coverage curve, not just the single chosen point. linear_api sits between the two on both axes, as it did on every split tested.

H. CROSS-MODEL REPLICATION CHECK: GPT-5.6 TERRA

Everything above uses a single model. The calibration failure mode documented in Sections V-A–V-E is argued as a general risk for selective prediction on LLM confidence scores, not a Claude-specific quirk—but that argument is stronger with direct cross-model evidence than without it. This section reruns the identical protocol—own from-scratch calibration on the $n = 120$ split, the corrected LCB selection rule, a single-use $n = 300$ holdout—against a second model, GPT-5.6 Terra (OpenAI's positional equivalent to Claude Sonnet 5: neither the top-tier nor the budget model in its family), on the linear and agentic flows (linear_gpt/agentic_gpt). direct_control and the self-consistency ablation (Section VI) were not rerun for this model; see Limitations. Every other element of the comparison is held fixed across both models: the same SOP text, the same six tools, the same case splits, and the same statistical pipeline—GPT's band is never seeded from or compared against Claude's frozen band, since reusing it would itself be exactly the kind of methodology leakage this benchmark's frozen-manifest guard exists to prevent.

Calibration ($n = 120$). linear_gpt selected a band organically: half-width 0.30, band (0.2, 0.8), LCB accuracy 0.884 (mean 0.931), 39/120 decided cases—comfortably clearing both the accuracy target and the minimum-decided-case floor. agentic_gpt did not: no candidate band cleared both bars. The closest, half-width 0.30, landed at LCB 0.849 (just under the 0.85 target) with 34 decided cases; the next step up, half-width 0.35, reached a perfect LCB 1.0 but only 21 decided cases, and calibrate_band()'s minimum-support guard correctly rejected it as a suspiciously perfect score from too few cases rather than accepting it. Selection fell through to the widest candidate band, (0.1, 0.9), itself supported by only 12 decided cases in calibration—a genuinely thin, unforced result of applying the identical corrected methodology to a




FIGURE 2. Full risk-coverage curves for all five arms that decide cases (direct_control decided none and is omitted), with each Claude/classical-ML arm's actual reported operating point marked (Section V-H's GPT arms omitted from the marked-point overlay for legibility; see Table 7 for their operating points). The vertical axis plots selective accuracy $acc(\tau) = 1 - R(\tau)$ ((2)), so a higher curve is a lower-risk curve; AUROC integrates the risk axis, so lower is better. classical_ml's curve dominates every LLM arm from either model across nearly the entire coverage range, not merely at the specific chosen operating points—AUROC 0.0343 (classical_ml) vs. 0.1080 (linear_api) vs. 0.1117 (agentic_api) vs. 0.1850 (linear_gpt) vs. 0.2046 (agentic_gpt).

new model, not a parameter chosen to produce a particular outcome.

TABLE 7. holdout_v2 results for the GPT-5.6 Terra arms ($n = 300$).

Arm	Band	Cov.	Acc.	Prec.	Rec.	F1	Cov-wt.
linear_gpt	(.2,.8)	24.3%	0.918	0.917	1.000	0.957	22.3%
agentic_gpt	(.1,.9)	6.0%	1.000	1.000	1.000	1.000	6.0%

TABLE 8. Bootstrap 95% CIs, GPT-5.6 Terra arms (10,000 resamples).

Arm	Accuracy	Coverage	Cov-wt.
linear_gpt	0.918 [0.849, 0.974]	0.244 [0.197, 0.293]	0.224 [0.180, 0.273]
agentic_gpt	1.000 [1.000, 1.000]	0.060 [0.033, 0.087]	0.060 [0.033, 0.087]

agentic_gpt's calibration-time caution held up faithfully on completely fresh data: it decided only 18 of 300 holdout cases (6.0% coverage) but got all 18 right (accuracy, precision, recall, and F1 all 1.000). That is the most extreme version of the high-precision/low-coverage pattern seen anywhere in this paper—more extreme than agentic_api's own already-conservative 17.0%-coverage, 96.1%-accuracy result—and it traces directly back to calibration: a band whose support was 12 decided cases produced a holdout behavior that is, if anything, *more* conservative still. This is not evidence the

LCB guard is too strict; the fact that a band selected on thin support turned out to generalize this narrowly, rather than surprising the model with a burst of wrong high-confidence calls, is closer to evidence the guard's caution was well-founded. linear_gpt lands at 24.3% coverage and 91.8% accuracy—close to linear_api's own 24.7%-coverage cell at the same band in Table 4, but at meaningfully lower accuracy (0.918 vs. 0.959).

Read against the full risk-coverage curve (Fig. 2, now showing all five arms): classical_ml continues to dominate every LLM arm from either model across nearly the entire coverage range. Area under the risk-coverage curve (AUROC) for the GPT arms is 0.1850 (linear_gpt) and 0.2046 (agentic_gpt)—both worse (higher) than either Claude arm (0.1080, 0.1117) and far worse than classical_ml (0.0343). At every band swept in Table 4's style, linear_gpt is consistently *more willing to commit* than linear_api at the same band width (e.g. 59.7% vs. 50.7% coverage at (0.35, 0.65); 11.0% vs. 6.0% at (0.10, 0.90)) but consistently *less accurate* among the cases it does decide—evidence GPT-5.6 Terra's probability outputs are less sharply polarized toward the extremes than Claude's, a genuinely different calibration shape between the two models, not merely a shifted threshold. Pairwise McNemar's tests between the matched flows (linear_api vs. linear_gpt, $n = 68$ both-decided; agentic_api vs. agentic_gpt, $n = 17$; linear_gpt vs. agentic_gpt, $n = 18$) all return zero discordant pairs ($p = 1.0$)—uninformative given how few cases any two

of these arms both decide, consistent with Section V-F's point that these tests are underpowered by construction here and carry little evidentiary weight next to the bootstrap CIs and risk-coverage curves above.

Faithfulness. Both faithfulness methods (Section IV's methodology, applied identically) were rerun on the GPT arms. **Method 1** (deterministic numeric cross-referencing) shows both GPT arms at $\sim 99.6\%$ verified rate (linear_gpt: 11,868/11,912 claimed numbers verified; agentic_gpt: 10,624/10,671)—marginally *higher* than either Claude arm's 98.7–99.3%. **Method 2** (LLM-judge semantic check, 25-case samples with two and three judge-response parse failures excluded for linear_gpt and agentic_gpt respectively, judged by claude-haiku-4-5 regardless of which model produced the case under review, to avoid a self-grading conflict of interest) tells a different story: linear_gpt shows misattributions in 11/23 judged cases (47.8%), agentic_gpt in 8/22 (36.4%)—both higher than linear_api's 6/24 (25.0%) and agentic_gpt's rate roughly matching agentic_api's 8/24 (33.3%). This divergence between the two methods—GPT's numeric citations are marginally *more* reliable, but its semantic attribution is judged *less* reliable, particularly for linear_gpt—is exactly why this paper uses two independent faithfulness checks rather than one: a method that only catches fabricated numbers would have reported GPT as the more faithful model outright, which is not the full picture.

Interpretation. The classical-ML-dominance finding generalizes: under a second model, run through the identical protocol, the same practical conclusion holds—classical_ml remains the strongest choice for resolving the queue in practice, by a wide margin, on both coverage-weighted correctness and the full risk-coverage curve. The calibration-discretization risk itself generalizes in spirit but not in literal mechanism: agentic_api's failure (Section V-D) was a band picked confidently on a misleadingly skewed small sample that then broke on holdout; agentic_gpt's calibration instead correctly detected its own thin support and degraded to maximal conservatism, which then *held* on holdout. Both are the same underlying story—a 120-case calibration set is marginal support for calibrating a threshold against LLM-generated probabilities, and what that marginality produces depends on the specific model's output distribution—but they are different failure shapes, which is itself useful evidence that the fix (LCB selection plus a minimum-support floor) is doing real, model-agnostic work rather than being tuned to one model's specific discretization pattern.

VI. ABLATION: SELF-CONSISTENCY

Self-consistency (implemented as `resolve_with_second_opinion` in `src/selective_prediction.py`) draws a second independent probability estimate whenever the first lands inside the escalate band, and nudges the estimate past the band edge on agreement—intended to convert genuine near-misses into confident decisions without forcing a guess on true toss-ups. To test whether this mechanism actually changes outcomes, linear_api and agentic_api were re-run on the iden-

tical 300 holdout_v2 cases with self-consistency disabled (--override-self-consistency false), under the same frozen bands used for the primary report, (0.35, 0.65) and (0.25, 0.75) respectively—an architecture-vs-architecture comparison on a fixed decision threshold, not a re-tune.

TABLE 9. Self-consistency on vs. off, holdout_v2 ($n = 300$).

Arm	Condition	Coverage	Accuracy
linear_api	on (primary)	50.7%	0.914
linear_api	off (ablation)	51.7%	0.903
agentic_api	on (primary)	17.0%	0.961
agentic_api	off (ablation)	16.3%	0.980

McNemar's test, restricted to cases decided under both conditions: linear_api ($n = 135$ both-decided): 0 discordant pairs, $p = 1.0$. agentic_api ($n = 40$ both-decided): 0 discordant pairs, $p = 1.0$.

Result: no detectable effect. Neither coverage nor accuracy moves by more than calibration-split single-run stochasticity (Section V-B: ± 1.7 points for both LLM arms, measured via repeat runs on the $n = 120$ calibration set—the closest available estimate of run-to-run noise, though not measured directly on holdout_v2), and on every case decided under both conditions, self-consistency changed the correctness of exactly zero of them. As implemented, self-consistency is a mechanistically plausible but empirically inert intervention for this task and model at this operating point—a negative result, reported here for completeness. It also confirms the accuracy differences reported in Section V are attributable to the calibration-band fix, not to self-consistency's presence.

Incidental finding: self-consistency's design intent (spend extra calls only on genuinely ambiguous cases) is corroborated by cost: turning it off saved \$10.14 (linear_api) and \$6.85 (agentic_api) across 300 cases each—the mechanism is triggering roughly as often as expected given the bands, even though it is not moving the accuracy needle at this operating point.

VII. COST AND LATENCY

TABLE 10. Mean cost and latency per case.

Arm	Mean cost/case	Mean latency/case
direct_control	\$0.047	29.7s
linear_api	\$0.130	64.3s
agentic_api	\$0.132	57.8s
linear_gpt	\$0.053	21.4s
agentic_gpt	\$0.160	53.8s
classical_ml	$\sim \$0$ (local)	~ 0 s

Full holdout_v2 run ($n = 300$, 3 Claude API arms): \$92.84, zero errors across all 900 case runs (one transient truncation retried successfully under the unchanged frozen configuration, consistent with a 0.33% base rate for this failure

mode observed across both holdout rounds). Self-consistency ablation (Section VI, both Claude arms, same 300 cases, self-consistency off): \$61.77, confirming the mechanism's cost is concentrated on the intended ambiguous-case subset. The two GPT-5.6 Terra holdout_v2 runs (Section V-H) cost \$63.88 combined ($n = 300$ each, zero errors across 600 case runs); their $n = 120$ calibration runs cost \$26.06 combined. `linear_gpt` is `linear_api`'s cheapest and fastest analogue in this study—2.5x cheaper, 3x faster—though at lower coverage-weighted correctness (Table 7); `agentic_gpt` costs slightly more than `agentic_api` per case despite a real 47% cache-hit rate on input tokens, an unconfirmed cost difference worth a closer look before treating it as settled (docs/methodology_log.md, 2026-08-28 entry).

VIII. DISCUSSION

The calibration-band failure documented in Section V is not specific to fraud or to this benchmark. Any pipeline that (a) elicits a numeric confidence score from an LLM and (b) calibrates a decision threshold against a modestly sized labeled set should expect the same failure mode: LLM probability outputs are not the smooth, continuous quantities classical calibration methods assume—they cluster on round, prompt-anchored values (Section V-D), and a threshold-selection procedure that optimizes a point estimate of accuracy, rather than a variance-aware lower bound, can mistake small-sample class imbalance within one of those clusters for genuine signal. The fix used here—select on an LCB criterion (Section V-E) rather than a point estimate, and enforce a minimum decided-case count per candidate band—is a general prescription for this class of problem, not a fraud-specific patch, and Section V-B's stochasticity check (± 1.7 pt run-to-run noise, far smaller than the ~ 10 -point holdout regression this failure produced) is what rules out ordinary generation randomness as the explanation and leaves the selection procedure as the culprit.

This failure mode also has a methodological implication beyond its fix: selective prediction changes which cross-arm comparison is meaningful. A single point-accuracy number is not a fair comparison between arms that escalate at very different rates—`agentic_api` at 17% coverage and `classical_ml` at 85% coverage are not competing on the same axis, and Table 4's per-band comparison and Section V-F's risk-coverage curves (Fig. 2) are the correct way to read the results: `classical_ml` dominates the full curve (AURC 0.0343 vs. 0.108–0.112), not merely the specific operating points each arm happened to select.

Faithfulness and outcome accuracy are also separable results, not the same finding read two ways. Both LLM arms are highly faithful to the evidence they retrieve ($\sim 99\%$ deterministic verified rate, Section IV), and the residual genuine errors caught by the LLM-judge pass are concentrated in minor evidence-handling categories (miscounting, cross-record misattribution) rather than fabrication (docs/error_taxonomy.md). The most common source of an outright wrong disposition is error-taxonomy category 9—cases where every cited

fact is correct and the reasoning is coherent, but the case itself was genuinely hard (e.g. HOLD2_0104: a well-corroborated multi-signal false positive). That is ordinary task difficulty under real uncertainty, not evidence fabrication, and it is worth stating plainly: a system can be trustworthy in how it uses evidence while still being wrong on the cases that are actually hard, and conflating the two would misdiagnose where to invest further effort (better retrieval and grounding vs. better judgment on ambiguous cases).

Put together, the practical reading for a fraud-ops team evaluating this class of system is that classical ML remains the stronger default: it resolves far more of the queue at comparable or better accuracy (Table 3, coverage-weighted 77.7% vs. 16.3%) and dominates the full risk-coverage trade-off. The LLM agent's demonstrated value in this evaluation is a high-precision, low-coverage second opinion—its single most accurate operating point in this study (96.1%, `agentic_api` at its selected band) edges out classical ML's accuracy at the same band, but at less than a quarter of the coverage—not a general-purpose replacement for the classical model.

Section V-H's cross-model check is what turns this from "true of one model we happened to study" into a claim with at least some cross-model support. Run through the identical protocol, GPT-5.6 Terra reaches the same qualitative conclusion—classical ML dominates the full risk-coverage curve against every LLM arm from either model (AURC 0.0343 vs. 0.108–0.205)—and its `agentic` arm pushes the high-precision/low-coverage pattern further still (6.0% coverage at 100% accuracy), not away from it. That the two models' calibration procedures failed in genuinely different shapes—one a confidently wrong narrow band, the other a correctly-detected thin-support fallback that then held on fresh data—is itself informative: it suggests the underlying risk (a modestly sized calibration set meeting a specific model's probability-output distribution) is the general hazard, while its exact manifestation is model-dependent. A practitioner calibrating a selective-prediction threshold against a new LLM should expect the same class of risk this paper documents for Claude, not assume it away because the failure will not necessarily look identical.

A. ETHICAL CONSIDERATIONS

FraudOps-Bench uses only the public, already-anonymized IEEE-CIS Kaggle competition dataset [23]—no real customer personally identifiable information (PII), case notes, or institutional data were used or generated. The fraud-analyst SOP driving all LLM arms' reasoning was authored by the researchers for this benchmark and has not been validated against a practicing analyst's judgment (Section IX); it should not be read as a claim about real institutional practice. No fairness or demographic-parity audit was performed on any arm's dispositions, and IEEE-CIS provides no demographic attributes to audit against directly—a limitation worth flagging explicitly, since a fraud-decisioning system's false-positive and false-negative costs are not symmetric across a real customer population, even though that asym-

metry is outside what this dataset can measure. This work reports a research benchmark, not a deployed or deployment-ready system, and none of its dispositions were used to make a real decision about a real transaction or person.

IX. LIMITATIONS

This work has the following limitations:

- **Balanced queue, not an operational base rate.** Every split is sampled 50/50 fraud/non-fraud (Section III-A), a deliberate choice to keep enough positives in a 120-case calibration set and a 300-case holdout to estimate per-band accuracy at all—but a queue that no real fraud-review function sees. IEEE-CIS runs at 3.50% fraud overall and 7.96% within the device-bearing pool these cases are drawn from. Precision and coverage-weighted correctness are the metrics most sensitive to this: at operational prevalence every arm's precision would fall substantially, and the false-positive burden a fraud-ops team actually cares about is not measured here at all. The between-arm comparisons are unaffected, since all arms are scored on identical cases, but no absolute number in this paper is a deployment estimate.
- **Cases require an identity record.** Sampling is restricted to transactions carrying device information, 23.9% of IEEE-CIS (Section III-A), since the investigation tools have nothing to return otherwise. Findings therefore describe the evidence-available slice of the queue, not the cases where an agent would have least to work with—plausibly the harder ones.
- **Single dataset.** IEEE-CIS only. No cross-dataset generalization claim is made or supported by this work.
- **Two models, not a broad sweep.** Claude Sonnet 5 (primary) and GPT-5.6 Terra (Section V-H cross-model check) only. This is evidence the findings are not purely Claude-specific, not a claim they generalize broadly—other frontier providers (e.g. Gemini), open-weight models, and other tiers within either model family remain untested.
- **GPT-5.6 Terra was not evaluated as thoroughly as Claude Sonnet 5.** Section V-H's cross-model check covers calibration, the single-use holdout, and both faithfulness methods, but reruns neither the self-consistency ablation (Section VI) nor the repeat-run stochasticity check (Section V-B) for this model—self-consistency was left on (the default) throughout, and no run-to-run variance estimate exists for GPT-5.6 Terra to rule out ordinary generation stochasticity the way Section V-B does for Claude. `direct_control` was also not rerun for this model.
- **Self-authored SOP.** The fraud-analyst SOP was written by the researchers, not validated by a practicing fraud analyst. It may not reflect real institutional practice.
- **Ground truth.** IEEE-CIS isFraud labels are treated as ground truth without independent verification; known label-noise concerns in Kaggle competition data are not specifically audited here.

- **No human baseline.** This work reports agent-vs.-classical-ML comparisons only; no claim is made about matching or exceeding human analyst performance.
- **Calibration-set size relative to output discretization.** The root-cause finding in Section V is itself evidence that a 120-case calibration set is marginal for this class of method when the underlying model's probability outputs are discretized. Section V-H's `agentic_gpt` result—a band supported by only 12 decided calibration cases—is a second, independent illustration of the same underlying marginality, not a coincidence specific to one model. The fix (LCB-based selection) mitigates but does not eliminate this risk; a larger calibration set would be a more robust long-term fix.
- **McNemar's tests underpowered.** Section V-F's paired significance tests have low power given how different the arms' coverage profiles are; the bootstrap CIs and risk-coverage curves carry more of the evidentiary weight in this work.

X. CONCLUSION AND FUTURE WORK

FraudOps-Bench asked whether an LLM agent can emulate a fraud analyst's investigation process rather than just its output label, and evaluated that question with a selective-prediction framing that lets a system decline to decide rather than forcing a guess. The headline empirical result is a genuine tradeoff, not a clean win for either approach: under a correctly calibrated band, the agentic arm's peak accuracy claim holds up on fresh, never-touched data (96.1%, its best result on any split), but it earns that accuracy by escalating 83% of the queue, while a classical gradient-boosted-tree baseline remains the stronger practical performer across the full risk-coverage tradeoff. The more transferable result, we think, is not either arm's number but the calibration failure mode documented in Section V: LLM confidence-score discretization interacting badly with a threshold selected on a small calibration set is a risk any selective-prediction pipeline built on LLM probability outputs should design around, independent of the underlying task. Section V-H's cross-model check on GPT-5.6 Terra is a first piece of direct evidence for that transferability claim, not just an assertion of it: the same protocol, applied to a second model, reaches the same practical conclusion (classical ML dominates the full risk-coverage curve against every LLM arm tested) and surfaces a second, differently-shaped calibration difficulty traceable to the same root cause—a modestly sized calibration set meeting a specific model's probability-output behavior.

Several deliberate scope cuts bound what this paper claims. No human fraud analyst was involved in this evaluation round—there is no human baseline, and the SOP driving the LLM arms' reasoning has not been validated by a practicing analyst (Section IX); both are natural next steps rather than oversights, and we do not claim this system matches or exceeds human analyst judgment. A second, structurally different dataset—BAF [21] is the strongest candidate, since it would extend the benchmark to account-opening fraud

rather than the card-not-present setting studied here, a genuinely different investigation shape—was identified but deliberately deferred rather than integrated under this round’s scope. Deeper architecture ablations beyond the single self-consistency check reported in Section VI (supervisor-node necessity, tool-count sensitivity, in light of SOP-Bench’s [8] finding that larger tool registries can hurt task success) remain open. Cross-model replication is no longer fully open—Section V-H reports one additional model—but remains partial: GPT-5.6 Terra was not carried through the self-consistency ablation or the repeat-run stochasticity check (Section IX), and providers beyond Claude and OpenAI (e.g. Gemini) and open-weight models are untested. We see the calibration-failure diagnosis-and-fix methodology, more than any specific accuracy number reported here, as the part of this work most worth other selective-prediction practitioners carrying forward—reinforced, now, by its holding up on a second model rather than resting on one.

DATA AND CODE AVAILABILITY

The IEEE-CIS Fraud Detection dataset is a public Kaggle competition dataset [23] and is not redistributed in this repository, consistent with the competition’s terms; researchers can obtain it directly from Kaggle. The processed, benchmark-specific case files derived from it—the dev/calibration/holdout_v1/holdout_v2 splits used for evaluation (data/processed/*.jsonl)—are included in the repository, since they contain only the small per-case evidence packets used by this benchmark’s tools, not the raw competition data.

All code (agent implementation, tool definitions, SOP text, calibration/evaluation/statistics pipelines, and the frozen-methodology manifest described in Section III-E) is available at github.com/pramegh-uikey/FraudOps-Bench. The per-arm holdout_v2 run outputs underlying Tables 3–10 and Fig. 2 (parsed dispositions, metrics, faithfulness scores, the statistics report, and the risk-coverage figure) are provided as supplementary material.

ACKNOWLEDGMENTS

The authors used Claude (Anthropic) [24] to draft portions of the Introduction, Related Work, Discussion, and Conclusion and Future Work sections, to identify and verify the prior-work citations in the References section, and to integrate existing experimental data (the Section V-B repeat-run variance table and the error-taxonomy findings referenced in Section VIII) into the manuscript text. All experimental design, code, data collection, and statistical analysis were performed by the authors; all AI-assisted text and citations were reviewed and verified by the authors against the underlying source data (docs/methodology_log.md, docs/error_taxonomy.md, and the per-arm run outputs provided as supplementary material) before inclusion.

REFERENCES

- [1] Association of Certified Fraud Examiners, *Occupational Fraud 2024: A Report to the Nations*. Austin, TX, USA: ACFE, 2024. [Online]. Available: <https://legacy.acfe.com/report-to-the-nations/2024/>
- [2] UK Finance, *Half Year Fraud Report 2025*. London, U.K.: UK Finance, 2025. [Online]. Available: <https://www.ukfinance.org.uk/policy-and-guidance/reports/half-year-fraud-report-2025>
- [3] P. Grover, J. Xu, J. Tittelfitz, A. Cheng, Z. Li, J. Zablocki, J. Liu, and H. Zhou, “Fraud dataset benchmark and applications,” *arXiv:2208.14417*, Aug. 2022.
- [4] X. Tan, Y. Ma, and X. Zhang, “Understanding structured financial data with LLMs: A case study on fraud detection,” *arXiv:2512.13040*, Dec. 2025.
- [5] S. Shuster, E. Zloof, A. Shabtai, and R. Puzis, “FAA framework: A large language model-based approach for credit card fraud investigations,” *arXiv:2506.11635*, Jun. 2025.
- [6] B. Wei, Y. S. Tay, H. Liu, J. Pan, K. Luo, Z. Zhu, and C. Jordan, “CORTEX: Collaborative LLM agents for high-stakes alert triage,” *arXiv:2510.00311*, Sep. 2025.
- [7] S. Jajodia, M. Sultana, S. Majumdar, A. Taylor, and G. Vandenbergh, “Before you hand over the wheel: Evaluating LLMs for security incident analysis,” *arXiv:2603.06422*, Mar. 2026.
- [8] S. Nandi *et al.*, “SOP-Bench: Complex industrial SOPs for evaluating LLM agents,” *arXiv:2506.08119*, Jun. 2025.
- [9] S. Yao, N. Shinn, P. Razavi, and K. Narasimhan, “ τ -bench: A benchmark for tool-agent-user interaction in real-world domains,” *arXiv:2406.12045*, Jun. 2024.
- [10] E. Levi and I. Kadar, “IntellAgent: A multi-agent framework for evaluating conversational AI systems,” *arXiv:2501.11067*, Jan. 2025.
- [11] S. G. Patil, H. Mao, F. Yan, C. C.-J. Ji, V. Suresh, I. Stoica, and J. E. Gonzalez, “The Berkeley function calling leaderboard (BFCL): From tool use to agentic evaluation of large language models,” in *Proc. 42nd Int. Conf. Machine Learning (ICML)*, PMLR vol. 267, 2025, pp. 48371–48392.
- [12] LangChain Inc., “LangGraph” [Software]. GitHub, 2024. [Online]. Available: <https://github.com/langchain-ai/langgraph> (accessed Sep. 7, 2026).
- [13] Q. Wu *et al.*, “AutoGen: Enabling next-gen LLM applications via multi-agent conversation framework,” *arXiv:2308.08155*, Aug. 2023.
- [14] crewAI Inc., “CrewAI” [Software]. GitHub, 2023. [Online]. Available: <https://github.com/crewAIInc/crewAI> (accessed Sep. 7, 2026).
- [15] C. K. Chow, “On optimum recognition error and reject tradeoff,” *IEEE Trans. Inf. Theory*, vol. 16, no. 1, pp. 41–46, Jan. 1970, doi: 10.1109/TIT.1970.1054406.
- [16] Y. Geifman and R. El-Yaniv, “Selective classification for deep neural networks,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2017, pp. 4878–4887.
- [17] K. Tian, E. Mitchell, A. Zhou, A. Sharma, R. Rafailov, H. Yao, C. Finn, and C. D. Manning, “Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback,” in *Proc. 2023 Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2023, pp. 5433–5442.
- [18] M. Xiong, Z. Hu, X. Lu, Y. Li, J. Fu, J. He, and B. Hooi, “Can LLMs express their uncertainty? An empirical evaluation of confidence elicitation in LLMs,” in *Proc. Int. Conf. Learning Representations (ICLR)*, 2024.
- [19] Y. Dai and Y. Wang, “Rescaling confidence: What scale design reveals about LLM metacognition,” *arXiv:2603.09309*, Mar. 2026.
- [20] M. M. Campos, A. Farinhas, C. Zerva, M. A. T. Figueiredo, and A. F. T. Martins, “Conformal prediction for natural language processing: A survey,” *arXiv:2405.01976*, May 2024.
- [21] S. Jesus, J. Pombal, D. Alves, A. Cruz, P. Saleiro, R. Ribeiro, J. Gama, and P. Bizarro, “Turning the tables: Biased, imbalanced, dynamic tabular datasets for ML evaluation,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS) Datasets and Benchmarks Track*, 2022.
- [22] E. R. Altman, J. Blanus, L. von Niederhäusern, B. Egressy, A. Anghel, and K. Atasu, “Realistic synthetic financial transactions for anti-money laundering models,” in *Proc. Adv. Neural Inf. Process. Syst. (NeurIPS) Datasets and Benchmarks Track*, 2023.
- [23] IEEE Computational Intelligence Society and Vesta Corporation, “IEEE-CIS Fraud Detection,” Kaggle competition, 2019. [Online]. Available: <https://www.kaggle.com/competitions/ieee-fraud-detection> (accessed Sep. 7, 2026).
- [24] Anthropic, “Claude” [Large language model]. San Francisco, CA, USA: Anthropic, 2026. [Online]. Available: <https://claude.ai> (accessed Sep. 7, 2026).



AYUSHI AMBILKAR received her degree in chemical engineering from the Indian Institute of Technology (IIT) Roorkee, India. She has approximately three years of industry experience as a data scientist, including work in the fraud detection and prevention domain at Accertify, a former subsidiary of American Express. Her research interests include applying large language model agents to fraud investigation workflows and selective-prediction evaluation methodology.



PRAMEGH UIKEY received his degree in computer science engineering from the Indian Institute of Technology (IIT) Roorkee, India. He has approximately three years of industry experience as a data scientist. His research interests include tool-augmented large language model agents, benchmark design, and statistical evaluation of machine learning systems.

...